

Neural Networks & Deep Learning

Assignment – 4

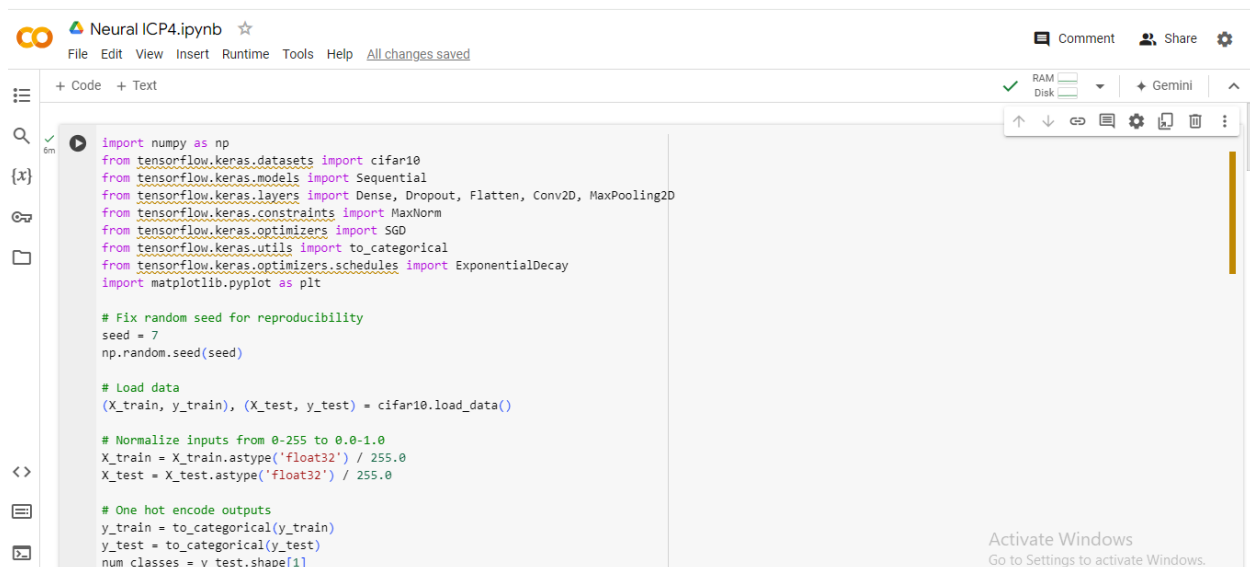
Name : Sravani Mannuru

Student Id : 700766162

Github Link: <https://github.com/sravs2031/Neural-Networks-ICP4.git>

Video Link: https://drive.google.com/file/d/1uDU3FUwF4NytDRDyofeXcREhflvN-tF1/view?usp=drive_link

Screenshots:



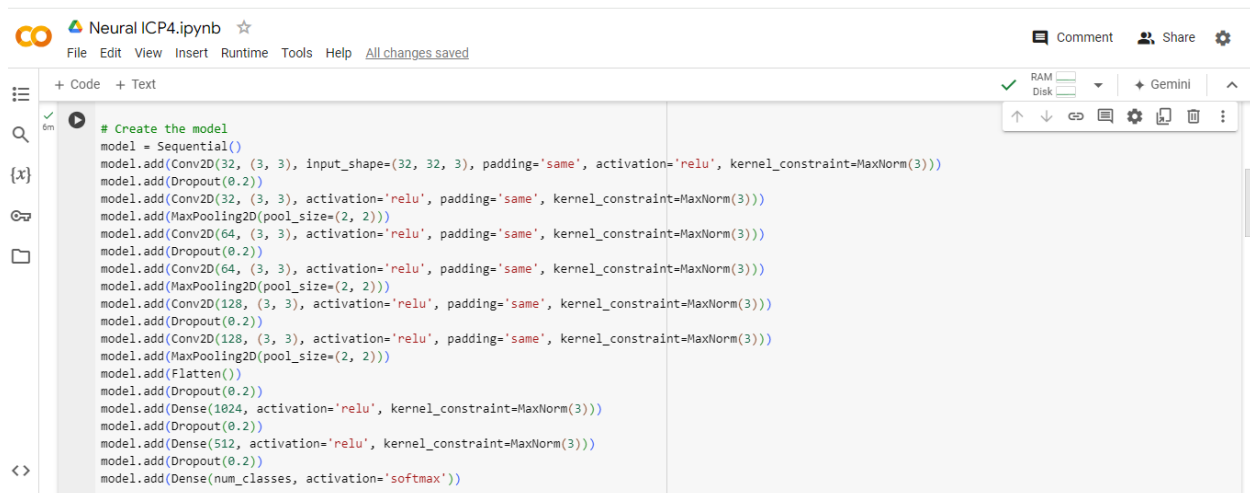
```
import numpy as np
from tensorflow.keras.datasets import cifar10
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout, Flatten, Conv2D, MaxPooling2D
from tensorflow.keras.constraints import MaxNorm
from tensorflow.keras.optimizers import SGD
from tensorflow.keras.utils import to_categorical
from tensorflow.keras.optimizers.schedules import ExponentialDecay
import matplotlib.pyplot as plt

# Fix random seed for reproducibility
seed = 7
np.random.seed(seed)

# Load data
(X_train, y_train), (X_test, y_test) = cifar10.load_data()

# Normalize inputs from 0-255 to 0.0-1.0
X_train = X_train.astype('float32') / 255.0
X_test = X_test.astype('float32') / 255.0

# One hot encode outputs
y_train = to_categorical(y_train)
y_test = to_categorical(y_test)
num_classes = y_test.shape[1]
```



```
# Create the model
model = Sequential()
model.add(Conv2D(32, (3, 3), input_shape=(32, 32, 3), padding='same', activation='relu', kernel_constraint=MaxNorm(3)))
model.add(Dropout(0.2))
model.add(Conv2D(32, (3, 3), activation='relu', padding='same', kernel_constraint=MaxNorm(3)))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Conv2D(64, (3, 3), activation='relu', padding='same', kernel_constraint=MaxNorm(3)))
model.add(Dropout(0.2))
model.add(Conv2D(64, (3, 3), activation='relu', padding='same', kernel_constraint=MaxNorm(3)))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Conv2D(128, (3, 3), activation='relu', padding='same', kernel_constraint=MaxNorm(3)))
model.add(Dropout(0.2))
model.add(Conv2D(128, (3, 3), activation='relu', padding='same', kernel_constraint=MaxNorm(3)))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Flatten())
model.add(Dropout(0.2))
model.add(Dense(1024, activation='relu', kernel_constraint=MaxNorm(3)))
model.add(Dropout(0.2))
model.add(Dense(512, activation='relu', kernel_constraint=MaxNorm(3)))
model.add(Dropout(0.2))
model.add(Dense(num_classes, activation='softmax'))
```

```
+ Code + Text
# Compile model
epochs = 5
lr_rate = 0.01
lr_schedule = ExponentialDecay(
    initial_learning_rate=lr_rate,
    decay_steps=epochs * len(X_train) // 32,
    decay_rate=0.1
)
sgd = SGD(learning_rate=lr_schedule, momentum=0.9, nesterov=False)
model.compile(loss='categorical_crossentropy', optimizer=sgd, metrics=['accuracy'])
print(model.summary())

# Fit the model
history = model.fit(X_train, y_train, validation_data=(X_test, y_test), epochs=epochs, batch_size=32)

# Final evaluation of the model
scores = model.evaluate(X_test, y_test, verbose=0)
print("Accuracy: %.2f%%" % (scores[1] * 100))

# Predict the first 4 images of the test data
predictions = model.predict(X_test[:4])
predicted_classes = np.argmax(predictions, axis=1)
actual_classes = np.argmax(y_test[:4], axis=1)
```

```
+ Code + Text
# Print the predictions and actual labels
print("Predicted classes: ", predicted_classes)
print("Actual classes: ", actual_classes)

# Plot the first 4 test images, predicted labels, and actual labels
fig, axes = plt.subplots(1, 4, figsize=(15, 3))
for i in range(4):
    axes[i].imshow(X_test[i])
    axes[i].set_title(f"Pred: {predicted_classes[i]}, Actual: {actual_classes[i]}")
    axes[i].axis('off')
plt.show()

# Visualize Loss and Accuracy
plt.figure(figsize=(12, 4))

# Plot Loss
plt.subplot(1, 2, 1)
plt.plot(history.history['loss'], label='train_loss')
plt.plot(history.history['val_loss'], label='val_loss')
plt.title('Model Loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(loc='upper right')

plt.xlabel('Epoch')
plt.legend(loc='upper right')
```

```
plt.xlabel('Epoch')
plt.legend(loc='upper right')

# Plot Accuracy
plt.subplot(1, 2, 2)
plt.plot(history.history['accuracy'], label='train_accuracy')
plt.plot(history.history['val_accuracy'], label='val_accuracy')
plt.title('Model Accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(loc='upper left')

plt.show()
```

Neural ICP4.ipynb

File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text

Download data from <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>
170498071/170498071 [=====] - 5s 0us/step
Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 32, 32, 32)	896
dropout (Dropout)	(None, 32, 32, 32)	0
conv2d_1 (Conv2D)	(None, 32, 32, 32)	9248
max_pooling2d (MaxPooling2D)	(None, 16, 16, 32)	0
conv2d_2 (Conv2D)	(None, 16, 16, 64)	18496
dropout_1 (Dropout)	(None, 16, 16, 64)	0
conv2d_3 (Conv2D)	(None, 16, 16, 64)	36928
max_pooling2d_1 (MaxPooling2D)	(None, 8, 8, 64)	0
conv2d_4 (Conv2D)	(None, 8, 8, 128)	73856
dropout_2 (Dropout)	(None, 8, 8, 128)	0

Activate Windows
Go to Settings to activate Windows.

Neural ICP4.ipynb

File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text

conv2d_3 (Conv2D)	(None, 16, 16, 64)	36928
max_pooling2d_1 (MaxPooling2D)	(None, 8, 8, 64)	0
conv2d_4 (Conv2D)	(None, 8, 8, 128)	73856
dropout_2 (Dropout)	(None, 8, 8, 128)	0
conv2d_5 (Conv2D)	(None, 8, 8, 128)	147584
max_pooling2d_2 (MaxPooling2D)	(None, 4, 4, 128)	0
flatten (Flatten)	(None, 2048)	0
dropout_3 (Dropout)	(None, 2048)	0
dense (Dense)	(None, 1024)	2098176
dropout_4 (Dropout)	(None, 1024)	0
dense_1 (Dense)	(None, 512)	524800
dropout_5 (Dropout)	(None, 512)	0
dense_2 (Dense)	(None, 10)	5130

Activate Windows
Go to Settings to activate Windows.

OUTPUT:

Neural ICP4.ipynb

File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text

```

=====
Total params: 2915114 (11.12 MB)
Trainable params: 2915114 (11.12 MB)
Non-trainable params: 0 (0.00 Byte)
=====
None
Epoch 1/5
1563/1563 [=====] - 419s 267ms/step - loss: 1.8524 - accuracy: 0.3143 - val_loss: 1.6051 - val_accuracy: 0.4140
Epoch 2/5
1563/1563 [=====] - 412s 264ms/step - loss: 1.4203 - accuracy: 0.4799 - val_loss: 1.2945 - val_accuracy: 0.5267
Epoch 3/5
1563/1563 [=====] - 410s 263ms/step - loss: 1.2258 - accuracy: 0.5584 - val_loss: 1.1829 - val_accuracy: 0.5712
Epoch 4/5
1563/1563 [=====] - 408s 261ms/step - loss: 1.0892 - accuracy: 0.6094 - val_loss: 1.0131 - val_accuracy: 0.6347
Epoch 5/5
1563/1563 [=====] - 409s 262ms/step - loss: 0.9947 - accuracy: 0.6463 - val_loss: 0.9915 - val_accuracy: 0.6500
Accuracy: 65.00%
1/1 [=====] - 0s 196ms/step
Predicted classes: [3 8 8 0]
Actual classes: [3 8 8 0]


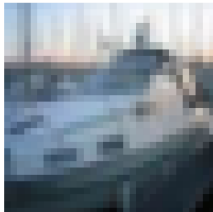
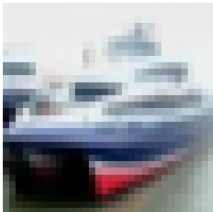
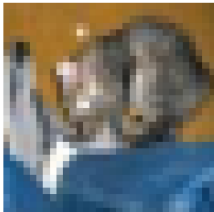
```

Neural ICP4.ipynb

File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text

Accuracy: 65.00%
1/1 [=====] - 0s 196ms/step
Predicted classes: [3 8 8 0]
Actual classes: [3 8 8 0]
Pred: 3, Actual: 3
Pred: 8, Actual: 8
Pred: 8, Actual: 8
Pred: 0, Actual: 0




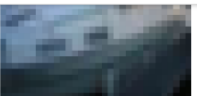


RAM Disk

+ Gemini

Neural ICP4.ipynb

File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text



Model Loss

Model Accuracy

train_loss

val_loss

train_accuracy

val_accuracy

